

The macroevolution of bipartite networks

Timothée Poisot ^{1,2,3} Daniel B. Stouffer ¹

1: Centre for Integrative Ecology, School of Biological Sciences, University of Canterbury, Christchurch, New Zealand; **2:** Université de Montréal, Département de Sciences Biologiques; **3:** Québec Centre for Biodiversity Sciences

This template uses the magic of makefiles, pandoc, and markdown, to make it easy to produce multiple documents from markdown, R markdown, or Julia markdown files. Just type `make` at the command line to see the different options.

Keywords: ecological networks - macroecoevolution - approximate Bayesian computation



This work is licensed under a Creative Commons Attribution 4.0 Unported License.

Correspondence to Timothée Poisot – timothee.poisot@umontreal.ca — Latest update on June 22, 2017

The extant structure and distribution of biodiversity is the outcome of macro-evolutionary processes, and the modelling of these processes has stimulated a large variety of approaches (???, ???). At their core, these approaches are all essentially birth-death processes, in that they model the rate of speciation and extinction to generate a prediction about both the temporal dynamics of species richness and its predicted current state. Surprisingly, these models tend to consider species as isolated entities; even though they share ancestry, they are not explicitly linked via inter-specific interactions. This fact is problematic from both an ecological (???) and evolutionary (???, ???, ???, ???) standpoint since it is widely accepted that interactions serve as an essential scaffold for biodiversity and its emergent properties such as community persistence or ecosystem function (???). After all, predators invariably require prey, hosts require parasites, flowering plants require pollinators, and so on.

Although modern macro-ecological models give an increasingly central role to interactions (???), such models are still unable to predict the structure of complex interacting communities (???). Nevertheless, there are two key observations upon which solutions to overcome this limitation can be devised. First, extant networks are decidedly non-random with regard to their structure, and their structure is equally non-random with regards to macro-evolutionary processes (???). Second, the structure of ecological networks is dynamic over evolutionary timescales (???). Both these points are strongly suggestive of perpetual and ongoing action of macro-evolutionary processes. It stands to reason then that models of macro-evolution with explicit consideration of species interactions will therefore provide an appropriate theoretical framework to understand how networks evolve. Notably, such a framework enables the estimation of how much of extant network structure originated through macro-evolution, as opposed to reflecting extant opportunities and constraints (???).

If one assumes that the conservatism of interactions across phylogenies can be explained by the fact that an incipient species inherits its ancestor's interactions upon speciation (???, ???), even a simple model with relatively few parameters can describe the possible evolutionary rules that shape a community's interaction network. Ideally, the parameters of any model such as this—no matter how simple or complex—ought to be calibrated against real-world evolutionary dynamics, similar to how the fossil and molecular record has been used to study species diversification (???). Unfortunately, the dearth of well-resolved, long-term time series of species interactions rules out such a comparison to temporal network dynamics. Therefore, we instead addressed the question of network macro-evolution here by using extant ecological networks to calibrate the end points of an interaction-centric birth-death simulation model under the assumption that the best-fitting models will provide insight into the network's likely evolutionary history. Among the variety of ecological networks types, bipartite ones are the most appropriate family to test this model: they have well partitioned interactions between guilds with no complex feedback loops, are present in a variety of systems and types of biological interactions, and there is a wealth of well-studied data available (???). Moreover, taxa from both guilds of a bipartite ecological network are usually tightly evolutionarily linked and require interactions to persist, making them ideal to elucidate evolutionary rules of community structure.

1 METHODS

1.1 Model description

We posit that four simple rules govern the evolution of networks. First, every network originally consists of just two species sharing a single interaction; for example, a plant and its herbivore. Second, a speciation event happens at the top level (*e.g.* the herbivore) with probability p , or at the bottom level with probability $1 - p$. Third, the incipient species starts with all interactions of its ancestor. Fourth, some of these interactions are lost with probability $\varepsilon(\lambda, k, c)$, which allows interactions—that are gained through speciation—to be lost either at a fixed rate λ or as a function of the incipient species' degree k . The c parameter modulates this relationship further by influencing whether high degree of an ancestor increases, or decreases, the probability of the incipient species losing interactions. We have used the following formulation for ε :

$$\varepsilon(\lambda, k, c) = \frac{1}{1 + \left(\frac{1}{\lambda} - 1\right) \times c^{k-1}}. \quad (1)$$

The resulting probability of interaction loss, and its consequences on degree, is shown in fig. 1. The values of ε belong to $]0; 1[$. Note that, because species are duplicated upon a speciation event, the network still grows over time. If an incipient species should lose all of its interactions, then it fails to establish.

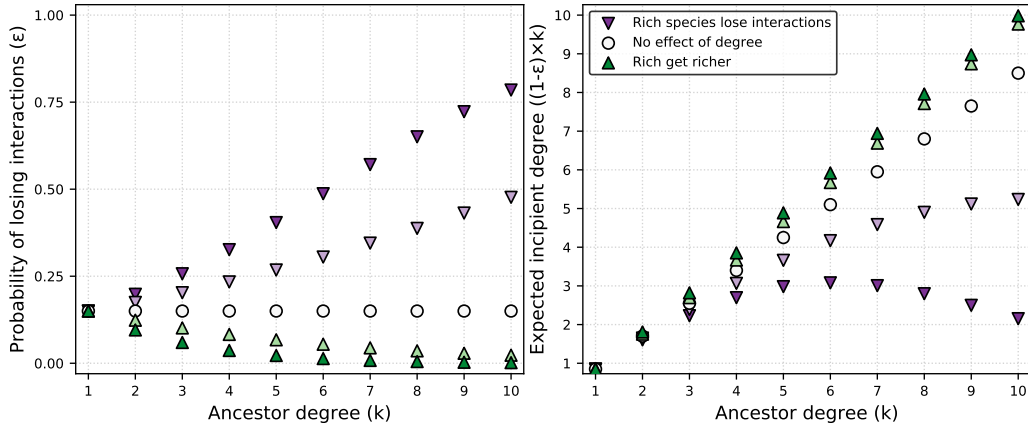


Figure 1 Left panel: probability that *each* interaction of the ancestor is lost by the incipient species during speciation. Right panel: xxx

The four previous steps are performed a fixed number of time – we impose an upper limit to the richness at each level, and when this limit is reached, the incipient species replaces one of the resident species at random. An equilibrium for the measures of network structure (see next section) is reached within 1000 timesteps. For all situations, we recorded the network after 5000 iterations.

1.2 Network measures

1.2.1 Connectance

Connectance, defined as the ratio of realized interactions on the total number of potential interactions, is one of the most common descriptor of network structure. In a bipartite network with T species at the top, and B at the bottom, having a total of L interactions, it is defined as $Co = L/(T \times B)$. Connectance has a lower bound, as the network cannot have fewer interactions than the number of species in its more speciose level – the minimal connectance is therefore $cm = \max(T, B)$. This makes the connectance of networks of different sizes difficult to compare, especially since bipartite networks tends to have a low connectance. For this reason, we used a corrected version of connectance, defined as

$$Co^* = \frac{L - cm}{T \times B - cm} . \quad (2)$$

This takes values between 0 (the network has the minimal number of interactions) and 1 (all species are connected), but is robust to variations in species richness.

1.2.2 Nestedness

We measured nestedness, using the η measure of Bastolla et al. (2009), which returns a global nestedness score based on the fact that interactions of relatively specialized species should be a subset of the interactions of more generalized ones. This measure is robust to changes in species richness, and returns values between 0 (not nested) to 1 (perfectly nested).

1.2.3 Modularity

We measured modularity using LP-BRIM (Liu & Murata 2009) (preliminary analyses revealed no qualitative impact of using other methods to optimize modularity). LP-BRIM returns values close to 1 when there are modules in the network, and values closer to 0 otherwise.

1.2.4 Motifs

Finally, we enumerated five four-species bipartite motifs (Baker et al. 2014). Bipartite motifs are possible conformations of four species spread across two levels, such as for example three consumers sharing one resource, or two consumers both exploiting resources, *etc.*. The five motifs we used are illustrated in fig. 2. Because the number of motifs obviously varies with species richness, we corrected it in the following way. For each level, we enumerated the sets of species with a degree allowing them to be part of the motif. Then we multiplied the number of sets for the top (t_x) and the bottom (b_x) level – this gives the number of possible combinations of species that *could* form a given motif. We then divided the count of observed motifs (m_x) by this product,

so that

$$m_x^* = \frac{m_c}{t_c \times b_x} \quad (3)$$

is the *proportion* of species that could form motif x which are actually in the correct conformation. This allows to compare the number of motifs between networks of different sizes.




Figure 2 Illustration of the five motifs used in this study. Motifs 21, 22, and 23 have the same number of species but different numbers of interactions; motifs 31 and 32 are flipped version of one another, and should help discriminate top-rich or bottom-rich communities.

1.2.5 Asymetry

Because empirical networks rarely have an equal number of species at both level, we measured asymetry as being the ratio of species at the top level on the species richness, so that

$$r = \frac{T}{T + B} . \quad (4)$$

1.3 Data selection

We used empirical data of plant-pollinator interactions (59 networks), plant-herbivore interactions (23 networks), phage-bacteria networks (38 interactions), plant-dispersers interactions (30 networks), and host-parasite interactions (121 networks). Pollination and seed-dispersal interactions come from the *WebOfLife* dataset (<http://mangal.io/data/dataset/7/>). Phage-bacteria (which are functionally equivalent to host-parasitoid) data are from (???). Host-parasite data (???) are from (???). Plant-herbivore data are from (???). Every network was “cleaned” in the following way. First, species with no interactions (if any) were removed. This yields networks in which all species have at least one interaction. Second, interactions strengths (if present) were removed since our model only requires information about the presence or absence of interactions.

1.4 Simulations

We conducted the following two numerical experiments. First, we conducted a systematic exploration of the model’s behaviour using evenly spaced parameter values. Each combination of parameters was simulated 1000 times. This allowed us to ensure that the model could return networks with all possible configurations, and that the output covered a range of network structures larger than what was observed in nature. Second, we sampled the parameter space uniformly, by drawing 10^5 parameter sets at random from within the aforementioned bounds. These outputs were used in the parameter selection experiment described below.

Each timestep in the simulation consists of three sub-steps. First, a level is picked at random: the top-level is picked with probability p , and the bottom-level is picked with probability $1 - p$. This is independent of the number of species at each level. Second, one species from the selected level is picked at random (all species within a level have equal chance of being picked), and duplicated (*i.e.* a novel species with the same interactions is added to the network). Each interaction of the incipient species is then *removed* with probability

In this formulation, k is the number of interactions of the incipient species, λ is the *basal* rate of interaction loss, and c is a parameter regulating whether species with more interactions tend to gain or lose interactions over time. Negative values of c imply that *rich get richer*, *i.e.* species with more interactions tend to conserve them more over speciation. The special case of $c = 0$ corresponds to no relationship between the degree of a species and its probability of losing or retaining an interaction over speciation.

We ran the model for 10^4 timesteps, for 10^5 random combinations of $< p, \lambda, c >$. Whenever either level has more than 10^2 species, some are deleted at random within this level. This ensure that the network is at most composed of 200 species. Preliminary analyses revealed that this threshold had no impact on the results presented as long as it was reasonably large (≥ 50).

1.5 Parameter selection

We used ABC (Approximate Bayesian Computation) to select the parameter values that yielded realistic networks by assessing how closely each replicate of the second numerical experiment resembles empirical communities. For each empirical network, its observed set of summary statistics was compared to each output of the stochastic model. The inverse of the Euclidean distance between the two arrays was recorded as the score of the parameter set. Because each empirical network is in practice a different optimization problem submitted to the ABC routine, and because ABC requires to set the rejection threshold on a per-problem basis, setting a global value was not meaningful. To circumvent this problem, we instead selected the posterior distribution as the 500 parameters sets that gave the best scores (*i.e.* above the 95th percentile). The distribution of distances (*i.e.* how well each point within the posterior distributions actually describes the empirical network) is kept to evaluate the global fit on a per-network basis.

2 RESULTS

Following our macro-evolutionary model, we repeated its four steps 10^4 times to generate a large ensemble of model networks whose structure we could compare to those of the empirical networks. We then compared these model-generated networks to a large collection of 271 bipartite ecological networks whose interactions encode seed dispersal, herbivory, parasitism, bacteriophagy or pollination (see *Methods*) using Approximate Bayesian Computation (ABC). When no analytical expression of a model's likelihood can be derived, ABC (???, ???) gives estimates of the posterior distributions of best-fit parameters (*i.e.* the most likely parameter values given the empirical data) by comparing a measure of distance between empirical observations and a model. Here, we define the distance between a simulated (*i*) and empirical (*j*) network as $d(\mathbf{v}_i, \mathbf{v}_j)$, where \mathbf{v} is an array of network structural properties, including connectance, modularity (???, ???), nestedness (Bastolla et al. 2009), and the distribution of different network motifs (???) (see *Methods*). For each network, the posterior distribution of best-fitting parameters is given by the set parameters of the closest 500 simulated models (to top 1% of the total).

!{posteriors}

We first observed that the posterior distribution of the parameters differs across interaction types (??). The probability of speciation at either level (p) is the least strongly selected, which suggests that mechanisms pertaining to the evolution of interactions have a stronger impact on extant network structure than does the distribution of speciation rates. We also encountered two situations for the distribution of the interaction rate λ : herbivory and pollination networks have higher values of this parameter, implying that herbivores and pollinators tend to retain the interactions of their ancestors more than other types of top-level organisms did (???, ???). All other types of networks were best described by low values of λ ; their interactions consequently appear to be more labile throughout the course of macro-evolution. Finally, all systems show a strong bias towards moderately high values of c ; this indicates that the effective probability of a species retaining its ancestor's interactions decreases with its ancestor's degree. That is, the generalism of species over time has an emergent upper bound, a fact that results in the very spectrum of high-degree and low-degree species that is ubiquitous empirically (???)

The optimal values of λ and c , however, are not independent since they ultimately affect the same process: the probability of the incipient species losing its ancestor's interactions. A more thorough understanding of the dynamics of interactions throughout evolution can therefore be obtained by examining these parameters' joint distribution. Doing so reveals two additional "states" for networks to occupy based on the results of our model (??); either c is close to 0 and λ is large or c is close to 1 and λ is low. Notably, different types of networks fall in a specific place within this continuum. Herbivory and pollination tend to have both low values of c and low to high values of λ —implying that the control on interaction persistence is at the community level—whereas parasitism networks have low values of λ and low-to-high values of c —implying that the control on interaction persistence is at the species level. The two remaining network types, seed dispersal and bacteriophagy, do not show a strong signal as to their position alongside this gradient.

!{parameters}

For each network, we next calculated the average distance to all its best matching simulation outputs, and used the z-score of this value to determine which type of networks was best predicted using our model (??). The best predicted networks were herbivory and pollination; this suggest that these networks have a particularly strong macro-evolutionary signal (???).

!{zscores}

Finally, we applied a classification tree to the parameter values describing each empirical network (??). The tree had a misclassification rate of 35.4%, meaning that knowing only the value of parameters λ and c , the correct type of ecological interaction can be estimated in around 65% of cases. The structure of tree also reveals that antagonistic and mutualistic interactions *do not* form different clusters (as opposed to what has been hypothesized before; ???), which contradicts the frequent assumption that different *consequences* of the interaction should imply different macro-evolutionary rules and trajectories (???).

!{tree}

Our results demonstrate that the structure of extant bipartite networks can be adequately reproduced by a speciation/extinction model that accounts for biotic interactions. The selection on parameters related to interaction diversification and persistence was stronger than on the parameter related to the rate of speciation, suggesting that the importance of biotic interactions in macro-evolution may have been understated. Our results also highlight that, while the evolutionary persistence of interactions is undeniably important in the macro-evolution of community structure, different type of ecological interactions respond in largely different ways. This offers a very stimulating possibility – namely, that because the mode of coevolution *within* the interaction between two species differ as a function of their ecological interactions (???), this can cascade up to the macro-evolutionary scale in the form of a signal of long-term interaction persistence.

REFERENCES

- Baker et al.** (2014). Species' roles in food webs show fidelity across a highly variable oak forest. *Ecography*. 38:130–9.
- Bastolla et al.** (2009). The architecture of mutualistic networks minimizes competition and increases biodiversity. *Nature*. 458:1018–20.
- Liu & Murata.** (2009). Community Detection in Large-Scale Bipartite Networks. *2009 IEEE/WIC/ACM International Joint Conference on Web Intelligence and Intelligent Agent Technology*. Institute of Electrical & Electronics Engineers (IEEE);